

```
# Import Libraries
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms, models
from torch.utils.data import DataLoader
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Device setup
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
# 1. Data Preprocessing
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
])

train_data = datasets.CIFAR10(root='./data', train=True, transform=transform, download=True)
test_data = datasets.CIFAR10(root='./data', train=False, transform=transform, download=True)

train_loader = DataLoader(train_data, batch_size=64, shuffle=True)
test_loader = DataLoader(test_data, batch_size=64, shuffle=False)
```

100%|██████████| 170M/170M [00:13<00:00, 12.3MB/s]

```
# 2. Load Pre-trained Model
model = models.resnet18(pretrained=True)
```

```
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated
  warnings.warn(
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|██████████| 44.7M/44.7M [00:00<00:00, 211MB/s]
```

```
# 3. Freeze feature extractor layers
for param in model.parameters():
    param.requires_grad = False
```

```
# 4. Modify final layer for CIFAR10 (10 classes)
num_features = model.fc.in_features
model.fc = nn.Linear(num_features, 10)
model = model.to(device)
```

```
# 5. Define Loss and Optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters(), lr=0.001)
```

```
# 6. Training Loop
train_losses, test_accuracies = [], []
for epoch in range(5): # small epochs for simplicity
    model.train()
    running_loss = 0.0
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    train_losses.append(running_loss / len(train_loader))

    # Evaluate accuracy
    model.eval()
    correct, total = 0, 0
    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
```

```

outputs = model(images)
_, predicted = torch.max(outputs, 1)
total += labels.size(0)
correct += (predicted == labels).sum().item()
acc = 100 * correct / total
test_accuracies.append(acc)
print(f"Epoch [{epoch+1}/5] - Loss: {train_losses[-1]:.4f} - Test Accuracy: {acc:.2f}%")

Epoch [1/5] - Loss: 0.8472 - Test Accuracy: 78.73%
Epoch [2/5] - Loss: 0.6233 - Test Accuracy: 79.74%
Epoch [3/5] - Loss: 0.5949 - Test Accuracy: 79.75%
Epoch [4/5] - Loss: 0.5790 - Test Accuracy: 80.36%
Epoch [5/5] - Loss: 0.5685 - Test Accuracy: 80.33%

```

```

# 7. Classification Report
y_true, y_pred = [], []
model.eval()
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(predicted.cpu().numpy())

print("\nClassification Report:\n")
print(classification_report(y_true, y_pred, target_names=train_data.classes))

```

Classification Report:

	precision	recall	f1-score	support
airplane	0.80	0.82	0.81	1000
automobile	0.94	0.83	0.88	1000
bird	0.82	0.68	0.74	1000
cat	0.73	0.66	0.69	1000
deer	0.71	0.81	0.76	1000
dog	0.78	0.77	0.77	1000
frog	0.86	0.81	0.83	1000
horse	0.75	0.88	0.81	1000
ship	0.84	0.88	0.86	1000
truck	0.84	0.91	0.88	1000
accuracy			0.80	10000
macro avg	0.81	0.80	0.80	10000
weighted avg	0.81	0.80	0.80	10000

```

# 8. Accuracy/Loss Graph
plt.figure()
plt.plot(train_losses, label='Training Loss')
plt.plot(test_accuracies, label='Test Accuracy')
plt.legend()
plt.title('Training Loss vs Test Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Value')
plt.show()

```

